

Measuring and Reporting Performance (Module 3 of Syllabus)

- *The administrator of a large data processing*
 - center may be interested in increasing *throughput*—the *total amount of work* done in a given time.
- we often want to relate the performance of two different computers, say, X and Y. The phrase "X is faster than Y" is used
 - to mean that the response time or execution time is lower on X than on Y for the given task.
 - "X is *n times faster than Y*" will mean
 - $n = \text{Execution Time}_y / \text{Execution Time}_x$

Why know about performance

- Purchasing Perspective:
 - Given a collection of machines, which has the
 - Best Performance?
 - Lowest Price?
 - Best Performance/Price?
- Design Perspective:
 - Faced with design options, which has the
 - Best Performance Improvement?
 - Lowest Cost?
 - Best Performance/Cost ?
- Both require
 - Basis for comparison
 - Metric for evaluation


Many possible definitions of performance

- Every computer vendor will select one that makes them look good. How do you make sense of conflicting claims?



Introducing the
2.20 GHz Pentium®4
Processor

Built with Intel's 0.13 micron
technology, the new 2.20
GHz Pentium®4 processor
delivers significant
performance gains.



Q: *Why do end users need a new performance metric?*

A: End users who rely only on megahertz as an indicator for performance do not have a complete picture of PC processor performance and may pay the price of missed expectations.

Measuring Performance

- Latency (response time, execution time)
 - Minimize time to wait for a computation
- Energy/Power consumption
- Throughput (tasks completed per unit time, bandwidth)
 - Maximize work done in a given interval
 - $= 1/\text{latency}$ when there is no overlap among tasks
 - $> 1/\text{latency}$ when there is
 - In real processors there is always overlap (pipelining)
- All are important :
 - Architecture – Latency is important,
 - Embedded system – Power consumption is important,
 - Network – Throughput is important

Some Definitions

- Performance is in units of things/unit time
 - E.g., Hamburgers/hour
 - Bigger is better
- If we are primarily concerned with response time
 - $\text{Performance}(x) = \frac{1}{\text{execution_time}(x)}$
- Relative performance: “X is N times faster than Y”

$$N = \frac{\text{Performance}(X)}{\text{Performance}(Y)} = \frac{\text{execution_time}(Y)}{\text{execution_time}(X)}$$

From previous slide

- Since execution time is the reciprocal of performance, the following relationship holds:

$$n = \frac{\text{Execution time}_Y}{\text{Execution time}_X} = \frac{\frac{1}{\text{Performance}_Y}}{\frac{1}{\text{Performance}_X}} = \frac{\text{Performance}_X}{\text{Performance}_Y}$$

Relative Performance

"X is n times faster than Y" means:

$$\frac{\text{Execution time}_Y}{\text{Execution time}_X} = n$$

"X is $m\%$ faster than Y" means:

$$\frac{\text{Execution time}_Y}{\text{Execution time}_X} \times 100\% = m$$

Two notions of performance

Plane	DC to Paris	Speed	Passengers	Throughput (pmph)
747	6.5 hours	610 mph	470	286,700
Concorde	3 hours	1350 mph	132	178,200

- Which has higher performance?
 - Depends on the **metric**
 - Time to do the task (Execution Time, Latency, Response Time)
 - Tasks per unit time (Throughput, Bandwidth)
 - Response time and throughput are often in opposition

Throughput

- "the throughput of X is 1.3 times higher than Y" signifies here
 - that the number of tasks completed per unit time on computer X is 1.3 times the number completed on Y

Execution Time

- execution time can be defined in different ways depending on what we count.
 - The most straightforward definition of time is called ***wall-clock time***.
- ***response time, or elapsed time, which is the latency to complete a task, including***
 - disk accesses, memory accesses, input/output activities, operating system overhead—everything

Measuring, Reporting, and Summarizing Performance

- With multiprogramming, the processor works on another program
- while waiting for I/O and may not necessarily minimize the elapsed time of one program.
- need a term to consider this activity. ***CPU time recognizes*** this distinction
 - the time the processor is computing, *not including* the time waiting for I/O or running other programs. (Clearly, the response time seen by the user is the elapsed time of the program, not the CPU time.)

Three Components of CPU Performance

$$\text{CPU time}_{x,p} = \text{Instructions executed}_p * \text{CPI}_{x,p} * \text{Clock cycle time}_x$$

↑
Cycles Per Instruction



CPU Performance

- The Fundamental Law

$$\text{CPU time} = \frac{\text{seconds}}{\text{program}} = \frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

- Three components of CPU performance:

- Instruction count
- CPI
- Clock cycle time

	Inst. Count	CPI	Clock
Program	X		
Compiler (Technology)	X	X	
Inst. Set Architecture	X	X	X
μArch (Organization)		X	X
Physical Design(Hardware)			X

CPI - Cycles per Instruction

Let F_i be the frequency of type i instructions in a program. Then,
Average CPI:

$$\text{CPI} = \frac{\text{Total Cycle}}{\text{Total Instruction Count}}$$
$$= \sum_{i=1}^n \text{CPI}_i \times F_i \quad \text{where} \quad F_i = \frac{\text{IC}_i}{\text{Instruction Count}}$$

$$\text{CPU time} = \text{Cycle time} \times \sum_{i=1}^n (\text{CPI}_i \times \text{IC}_i)$$

Example:


Instruction type	ALU	Load	Store	Branch
Frequency	43%	21%	12%	24%
Clock cycles	1	2	2	2

average CPI = $0.43 + 0.42 + 0.24 + 0.48 = 1.57$ cycles/instruction

CPI

- The average number of clock cycles per instruction, or **CPI**, is a function of the machine and program.
 - The CPI depends on the actual instructions appearing in the program—a floating-point intensive application might have a higher CPI than an integer-based program.
 - It also depends on the CPU implementation. For example, a Pentium can execute the same instructions as an older 80486, but faster.
- It is common to each instruction took one cycle, making $CPI = 1$.
 - The CPI can be >1 due to memory stalls and slow instructions.
 - The CPI can be <1 on machines that execute more than 1 instruction per cycle (superscalar).

Clock cycle time



- One “cycle” is the minimum time it takes the CPU to do any work.
 - The **clock cycle time** or clock period is just the length of a cycle.
 - The **clock rate**, or frequency, is the reciprocal of the cycle time.
- Generally, a higher frequency is better.
- Some examples illustrate some typical frequencies.
 - A 500MHz processor has a cycle time of 2ns.
 - A 2GHz (2000MHz) CPU has a cycle time of just 0.5ns (500ps).

Execution time, again

$$\text{CPU time}_{x,p} = \text{Instructions executed}_p * \text{CPI}_{x,p} * \text{Clock cycle time}_x$$

- The easiest way to remember this is match up the units:

$$\frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Clock cycles}}{\text{Instructions}} * \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Make things faster by making any component smaller!!

	Program	Compiler	ISA	Organization	Technology
Instruction Executed					
CPI					
Clock Cycle Time					

- Often easy to reduce one component by increasing another

Speedup Performance Laws

1) Amdahl's Law

(Module 3 of Syllabus)

- Quantitative Principles of Computer Design
 - Make common case fast
 - Improving frequent event rather than rare event improves overall performance
 - i.e. in adding 2 numbers in CPU,
 - performance improved by optimizing common case of no overflow rather than rare case of flowchart
- **To find how much performance can be improved by making, a frequent event faster – Amdahl's law is used to quantify the values of enhancement (/ non-enhancement)**

Amdahl's Law

- A quick way to find speedup from an enhancement, which depends on:
 - Fraction of the computation time in the original machine that can be converted to take advantage of the enhancement
 - Example : If 20 seconds of execution time of a program that takes 60 seconds in total can use an enhancement, the fraction is **Fraction_{enhanced}** = $20/60$ is always less than 1
 - How much faster the task would run if the enhanced mode were used for the entire program
 - If the enhanced mode **takes 10 seconds for some portion of program** that can **completely use enhanced mode** and the **original mode took 20 seconds for the same portion**, the improvement is **Speedup_{enhanced}** = $20/10$ always greater than 1

Compute Speedup – Amdahl's Law

Speedup is due to enhancement(E):



$$\text{Speedup (E)} = \frac{\text{Execution time without E (Before)}}{\text{Execution time with E (After)}}$$

Suppose that enhancement E accelerates a fraction F (**Fraction_{enhanced}**) of the task by a factor S(**Speedup_{enhanced}**), and the remainder of the task is unaffected, what is the **Execution time_{after}** and **Speedup(E)**?

Amdahl's Law

$$\text{Execution time}_{\text{after}} = \text{ExTime}_{\text{before}} \times \left[(1 - F) + \frac{F}{S} \right]$$

$$\text{Speedup}(E) = \frac{\text{ExTime}_{\text{before}}}{\text{ExTime}_{\text{after}}} = \frac{1}{\left[(1 - F) + \frac{F}{S} \right]}$$

(F= fraction of total time that is speeded up ie 20 seconds out of total of 60 seconds
= 20/60

20 secs is speeded up to 10 seconds, say)

S= 20/10

Amdahl's Law :Example Given in Previous slide

ExTime_{before} = 60 , **Fraction**_{enhanced} (F) = 20/60, **Speedup**_{enhanced} (S)= 20/10

$$\text{Execution time}_{\text{after}} = \text{ExTime}_{\text{before}} \times \left[(1 - F) + \frac{F}{S} \right] = 60 \times 5/6 = 50$$

$$\begin{aligned} \text{Execution time}_{\text{after}} &= 60 - 60 \times 20/60 + 60 \times 20/60 \times 10/20 \\ &= 60 - 20 + 10 = 40 + 10 = 50 \end{aligned}$$

ExecutionTimeNon Enhanced Portion + Execution time of Enhanced portion

$$\left[(1 - F) + \frac{F}{S} \right] = (1 - 20/60) + 20/60 \times 10/20 = 2/3 + 1/6 = 5/6$$

(1-F) >> much greater than F/S i.e 2/3 >> 1/6

$$\begin{aligned} \text{Speedup}(E) &= \frac{\text{ExTime}_{\text{before}}}{\text{ExTime}_{\text{after}}} = \frac{6}{5} \\ &= 1.2 \end{aligned}$$

Check 60/50 = 1.2

Corollary : We can't Speedup the task more than the reciprocal of 1 minus the fraction , that is the portion that is not speeded up , here (1/(1-20/60)=3/2=1.5)

Amdahl's Law – An Example

Q: Floating point instructions improved to run 2X;
but only 10% of execution time are FP ops. What is the execution time
and speedup after improvement?

Ans:

$$F = 0.1, S = 2$$

$$\text{ExTime}_{\text{after}} = \text{ExTime}_{\text{before}} \times [(1-0.1) + 0.1/2] = 0.95 \text{ ExTime}_{\text{before}}$$

$$\text{Speedup} = \frac{\text{ExTime}_{\text{before}}}{\text{ExTime}_{\text{after}}} = \frac{1}{0.95} = 1.053$$

Amdahl's Law for Multiple Processors

- **based on fixed workload or fixed problem size**
- computational workload W is fixed while the number of processors that can work on W can be increased.
- Denote the execution rate of i processors as R_i then in a relative comparison they can be simplified as and $R_1=1 \dots R_n=n$.
- The workload is also simplified. We assume that the workload consists of sequential work $W\alpha$ and n parallel work $1-\alpha W$, where α is between 0 and 1. More specifically, this workload can be written in a vector form as, $(\alpha, 0, \dots, 0, 1-\alpha)W$, or $W_1=\alpha W$, $W_n=(1-\alpha)W$, $W_i=0$ for all $i \neq 1, n$.

Amdahl's Law for Multiple Processors

- The execution time of the given work by n processors is then computed as,

$$T_n = \frac{W_1}{R_1} + \frac{W_n}{R_n}$$

Speedup of n processor system is defined using a ratio of execution time, i.e.,

$$S_n = \frac{T_1}{T_n}$$

Amdahl's Law for Multiple Processors

Substituting the execution time in relation to W gives
:

$$S_n = \frac{W/1}{\frac{\alpha W}{1} + \frac{(1-\alpha)W}{n}} = \frac{n}{1 + (n-1)\alpha}$$

$$S_n = \frac{n}{1 + (n-1)\alpha}$$

Equation called Amdahl's Law

If the number of processors is increased infinity, the speedup becomes

$$S_\infty = \frac{1}{\alpha}$$

This is sequential bottle neck of multiprocessor system

The **speedup** can **NOT** be increased to infinity even if the number of processors is increased to infinity.

Gustafson's Law

- For scaled up problems(problem scaled to match the computing power of the machine as number of processors is increased)
- workload is scaled up to maintain a fixed execution time as the **number of processors increases**,
- workload is scaled up on an n-node machine as

$$W' = \alpha W + (1 - \alpha)nW$$

Gustafson's Law

- Speedup for the scaled up workload is then,

$$S'_n = \frac{\text{Single Processor Execution Time}}{n - \text{Processor Execution Time}}$$

$$S'_n = \frac{(\alpha W + (1 - \alpha)nW) / 1}{\frac{\alpha W}{1} + \frac{(1 - \alpha)nW}{n}}$$

Simplifying above equation produces the Gustafson's law:

$$S'_n = \alpha + (1 - \alpha)n$$

the speedup increases linearly.

- What Gustafson's law says
 - true **parallel power** of a large multiprocessor system is only **achievable when a large parallel problem is applied**.

MIPS and MFLOPS

- **MIPS:** millions of instructions per second:
 - $\text{MIPS} = \text{Inst. count} / (\text{CPU time} * 10^6) = \text{Clock rate} / (\text{CPI} * 10^6)$
 - easy to understand and to market
 - inst. set dependent, cannot be used across machines.
 - program dependent
 - can vary inversely to performance! (why? read the book)
- **MFLOPS:** million of FP ops per second.
 - less compiler dependent than MIPS.
 - not all FP ops are implemented in h/w on all machines.
 - not all FP ops have same latencies.
 - normalized MFLOPS: uses an equivalence table to even out the various latencies of FP ops.