

Module-2  
CSEN 3104  
Lecture 20  
26/08/2019

Dr. Debranjan Sarkar

# SIMD Algorithms

Sorting on a mesh-  
connected  
parallel computer

# Parallel Sorting on mesh

- Sorting of  $N = n^2$  elements on an  $n \times n$  mesh-type processor array
- Architecture (show figure) is similar to Illiac IV with exceptions
  - No wraparound connections, i.e.,
  - PEs at the perimeter have 2 or 3 rather than 4 neighbours
  - This simplifies the array sorting algorithm
- Two time measures are required to estimate the time complexity of the algorithm:
  - Routing time ( $t_R$ ) to move one data item from a PE to one of its neighbours
  - Comparison time ( $t_C$ ) for one comparison step (conditional interchange on the contents of two registers in each PE)

# Parallel Sorting on mesh

- Concurrent data routing is allowed
- Upto N numbers of concurrent comparisons may be performed
- This means that a comparison-interchange step between two items in adjacent processors can be done in time  $2t_R + t_C$  (*route left, compare, and route right*)
- A number of these comparison-interchange steps may be performed concurrently in time  $(2t_R + t_C)$  if they are all between distinct, vertically adjacent processors
- A mixture of horizontal and vertical comparison-interchanges will require at least  $(4t_R + t_C)$  time unit

# Parallel Sorting on mesh

- The PEs may be indexed by a bijection from  $\{1, 2, \dots, n\} \times \{1, 2, \dots, n\}$  to  $\{0, 1, \dots, N - 1\}$ , where  $N = n^2$
- $N$  elements of a linearly ordered set are initially loaded in the  $N$  PEs
- Sorting problem is defined as the problem of moving the  $j^{\text{th}}$  smallest element to the processor indexed by  $j$  for all  $j = 0, 1, \dots, N - 1$
- Example:
- The elements ( $N=16, n = 4$ ) to be sorted are initially loaded in the  $4 \times 4$  array of PEs (Show Figure)
- Three ways of indexing the processors
  - *Row-major indexing*
  - *Shuffled row-major indexing*
  - *Snake-like row-major indexing*

# Parallel Sorting on mesh

- *Row-major indexing (Show diagram)*
- *Shuffled row-major indexing (Show diagram)*
  - Note that this indexing is obtained by shuffling the binary representation of the row-major index
  - For example, the row-major index 5 has the binary representation 0101
  - Shuffling the bits gives 0011 which is 3
  - In general, the shuffled binary number, say, "abcdefgh" is "aebfcgdh"
- *Snake-like row-major indexing (Show diagram)*
  - Obtained from the row-major indexing by reversing the ordering in even rows
- The choice of a particular indexing scheme depends upon how the sorted elements will be used
- We are interested in designing algorithms which minimize the time spent in routing and comparing

# Parallel Sorting on mesh

- For any index scheme, there are situations where the two elements initially loaded at the opposite corner PEs, have to be transposed during the sorting (Show Figure)
- This transposition needs at least  $4(n - 1)$  routing steps
- This implies that no algorithm can sort  $n^2$  elements in time less than  $O(n)$
- Thus an  $O(n)$  sorting algorithm is considered optimal on a mesh of  $n^2$  PEs
- We shall show one such optimal sorting algorithm on the mesh-connected PEs



# Odd-even Transposition Sort

- Different Sorting algorithms
  - Bubble Sort Computational Complexity:  $O(n^2)$  in average case
  - Merge Sort Computational Complexity:  $O(n \log n)$  in worst case
  - Quick Sort Computational Complexity:  $O(n \log n)$  in average case
- These algorithms are not easily parallelizable
  - Because the operations depend on the result of the previous operations
- Odd-even Transposition sort (or Brick Sort) is suitable for parallel computers and the time complexity is reduced to  $O(n)$
- Examples of Odd-Even Transposition Sort

# Review of Batcher's odd-even merge sort

- Sort the first half of a list, and sort the second half separately
- Sort the odd-indexed entries (first, third, fifth, ...) and the even-indexed entries (second, fourth, sixth, ...) separately
- Make only one more comparison-switch per pair of keys to completely sort the list
- List of numbers: 2 7 6 3 9 4 1 8
- We wish to sort it from least to greatest
- If we sort the first and second halves separately we obtain: 2 3 6 7 / 1 4 8 9
- Sorting the odd-indexed keys (2, 6, 1, 8) we get (1 2 6 8)
- Sorting the even-indexed keys (3, 7, 4, 9) we get (3 4 7 9)
- Leaving them in odd and even places respectively yields: 1 3 2 4 6 7 8 9
- This list is now almost sorted
- Doing a comparison switch between the keys in positions (2 and 3), (4 and 5) and (6 and 7) will finish the sort

# Review of Batcher's odd-even merge sort

- Normally, the length of the list is a power of 2 (Here  $2^3 = 8$ )
- Two sorted sequences are loaded on a set of linearly connected PEs (Show Figure)
- In the first stage, the odd-indexed elements are placed in the left and then the even-indexed elements are placed in the right. This is basically unshuffle (or inverse shuffle) operation
- In the second stage, the odd sequences and the even sequences are merged
- The third stage is basically a perfect shuffle operation
- The fourth and final stage is a comparison-interchange operation of even-indexed elements with the next element
- Note that the perfect shuffle can be achieved by using the triangular interchange pattern (show figure)
- Similarly, an inverted triangular interchange pattern will do the unshuffle.
- The double-headed arrows indicate interchanges

Thank you