



# Deadlock

# Operating Systems

- Objectives
  - describe deadlock, and forms of prevention, avoidance, detection, and recovery

# Contents

1. What is Deadlock?
2. Dealing with Deadlock
3. Deadlock Prevention
4. Deadlock Avoidance
5. Deadlock Detection
6. Deadlock Recovery



# 1. What is Deadlock?

- An example from US Kansas law:
  - “When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.”

# In Picture Form:



# 1.1. System Deadlock

- A process must request a resource before using it, and must release the resource after finishing with it.
- A set of processes is in a *deadlock state* when every process in the set is waiting for a resource that can only be released by another process in the set.



## 1.2. Necessary Conditions for Deadlock

- Mutual Exclusion
  - at least one resource must be held in non-shareable mode
- Hold and Wait
  - a process is holding a resource and waiting for others

*continued*

- No Preemption

- only the process can release its held resource



- Circular Wait

- $\{P_0, P_1, \dots, P_n\}$
  - $P_i$  is waiting for the resource held by  $P_{i+1}$ ;  
 $P_n$  is waiting for the resource held by  $P_0$



# 1.3. Resource Allocation Graph

- A set of processes  $\{P_0, P_1, \dots\}$
- A set of resource types  $\{R_1, R_2, \dots\}$ , together with instances of those types.




$R_k$

# Edge Notation

- $P_i \rightarrow R_j$ 
  - process  $i$  has requested an instance of resource  $j$
  - called a *request edge*
- $R_j \rightarrow P_i$ 
  - an instance of resource  $j$  has been assigned to process  $i$
  - called an *assignment edge*

# Example Graph







# Finding a Deadlock

- If the graph has no cycles then there are no deadlock processes.
- If there is a cycle, then there *may* be a deadlock.

# Graph with a Deadlock



# Graph without a Deadlock






## 2. Dealing with Deadlocks

- Stop a deadlock ever occurring
  - *deadlock prevention*
    - disallow at least one of the necessary conditions
  - *deadlock avoidance*
    - see a deadlock coming and alter the process/resource allocation strategy

*continued*

- 
- *Deadlock detection and recovery*
  - Ignore the problem
    - done by most OSes, including UNIX
    - cheap solution
    - infrequent, manual reboots may be acceptable



# 3. Deadlock Prevention

- Eliminate one (or more) of:
  - mutual exclusion
  - hold and wait
  - no preemption (i.e. *have* preemption)
  - circular wait



# 3.1. Eliminate Mutual Exclusion

- Shared resources do not require mutual exclusion
  - e.g. read-only files
- But some resources cannot be shared (at the same time)
  - e.g. printers

## 3.2. Eliminate Hold & Wait

- One approach requires that each process be allocated *all* of its resources before it begins executing
  - eliminates the wait possibility
- Alternatively, only allow a process to request resources when it currently has none
  - eliminates the hold possibility



## 3.3. Eliminate “No Preemption”

- Make a process automatically release its current resources if it cannot obtain all the ones it wants
  - restart the process when it can obtain everything
- Alternatively, the desired resources can be preempted from other waiting processes

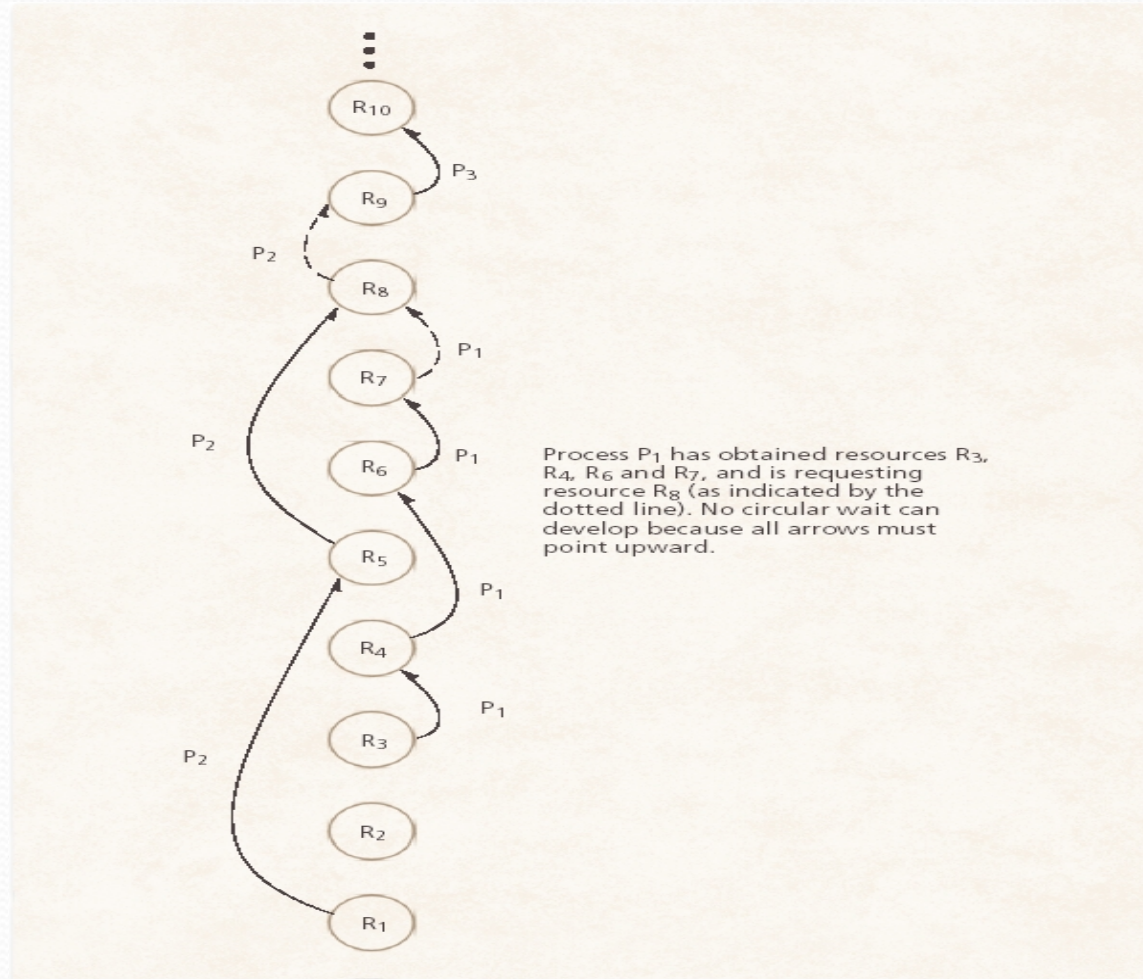


## 3.4. Eliminate Circular Wait

- Impose a total ordering on all the resource types, and force each process to request resources in increasing order.
- *Another approach*: require a process to release larger numbered resources when it obtains a smaller numbered resource.

# 3.4 Eliminate Circular Wait

Havender's linear ordering of resources to prevent deadlock





# 4. Deadlock Avoidance

- In deadlock avoidance, the necessary conditions are untouched.
- Instead, extra information about resources is used by the OS to do better forward planning of process/resource allocation
  - indirectly avoids circular wait




## 4.1. Safe States

- An OS is in a *safe state* if there is a *safe sequence* of process executions  $\langle P_1, P_2, \dots, P_n \rangle$ .
- In a safe sequence, each  $P_i$  can satisfy its resource requests by using the currently available resources *and* (if necessary) the resources held by  $P_j$  ( $j < i$ )
  - only when  $P_j$  has finished

# Safe State Implications

- A safe state cannot lead to deadlock.
- An unsafe state *may* lead to deadlock.
- Deadlock is avoided by always keeping the system in a safe state
  - this may reduce resource utilization





# Example

## 1

- Max no. of resources: 12 tape drives

- |       | <u>Max needs</u> | <u>Current Allocation</u> |
|-------|------------------|---------------------------|
| $P_0$ | 10               | 5                         |
| $P_1$ | 4                | 2                         |
| $P_2$ | 9                | 2                         |

- Currently, there are 3 free tape drives
- The OS is in a *safe* state, since  $\langle P_1, P_0, P_2 \rangle$  is a safe sequence.



# Example 2

- Same as last slide, but  $P_2$  now has 3 tape drives allocated currently.

	<u>Max needs</u>	<u>Current Allocation</u>
$P_0$	10	5
$P_1$	4	2
$P_2$	9	<u>3</u>

- The OS is in an *unsafe* state.

## 4.2. Using Resource Allocation Graphs

- Assume a resource type only has one instance.
- Add a *claim edge*:
  - $P_i \rightarrow R_j$
  - process  $P_i$  may request a resource  $R_j$  in the future
  - drawn as a dashed line


*continued*




- When the resource is actually requested, the claim edge is changed to a request edge.
- When an assignment is released, the assignment edge is changed back to a claim edge.

*continued*



- 
- All resources must be claimed before system start-up.
  - An unsafe state is caused by a cycle in the resource allocation graph.

# Example



→  
 $R_2$  allocation to  $P_2$   
creates an unsafe state

## 4.3. Banker's Algorithm

- Assume that:
  - a resource can have multiple instances
  - the OS has  $N$  processes,  $M$  resource types
- Initially, each process must declare the maximum no. of resources it will need.
- Calculate a safe sequence if possible.



# Banker Data Structures

- Available[M]
  - no. of available resource instances for each resource type
  - e.g. Available[j] == k means K R<sub>j</sub>'s
- Max[N][M]
  - max demand of each process
  - e.g. max[i][j] == k means P<sub>i</sub> wants k R<sub>j</sub>'s

*continued*

- `Work[M]`
  - no. of resource instances available for work (by all processes)
  - e.g. `Work[j] == k` means  $K$   $R_j$ 's are available
- `Finish[N]`
  - record of finished processes
  - e.g.  $P_i$  is finished if `Finish[i] == true`

*continued*



- Allocation[N][M]
  - no. of resource instances allocated to each process
  - e.g. Allocation[i][j] == k  
means  $P_i$  currently has k  $R_j$ 's
- Need[N][M]
  - no. of resource instances still needed by each process
  - e.g. Need[i][j] == k  
means  $P_i$  still needs k  $R_j$ 's
  - Need[i][j] == Max[i][j] - Allocation[i][j]

*continued*

- Request[N][M]
  - no. of resource instances currently requested by each process
  - e.g. Request[i][j] == k means  $P_i$  has requested k  $R_j$ 's



# Vectors

shorthand for  
referring to the  
2D data structures

- `Allocation[i]`
  - resources currently allocated to  $P_i$
- `Need[i]`
  - resources still needed by  $P_i$
- `Request[i]`
  - resources currently requested by  $P_i$

# The Safety Algorithm

- 1 Vector Copy:  $Work := Available$ ;  $Finish := false$
- 2 Find  $i$  such that  $P_i$  hasn't finished but could:  
 $Finish[i] == false$   
 $Need[i] \leq Work$   
If no suitable  $i$ , go to step 4.
- 3 Assume  $P_i$  completes:  
 $Work := Work + Allocation[i]$   
 $Finish[i] := true$   
go to step 2
- 4 If for all  $i$   $Finish[i] == true$  then *Safe-State*



# Safety Example

- | <u>Resource Type</u> | <u>Instances</u> |
|----------------------|------------------|
| A                    | 10               |
| B                    | 5                |
| C                    | 7                |

*continued*

			<u>Allocation</u>			<u>Max</u>			<u>Available</u>			<u>Need</u>		
A	B	C				A	B	C	A	B	C	A	B	C
P <sub>0</sub>			0	1	0	7	5	3	3	3	2	7	4	3
P <sub>1</sub>			2	0	0	3	2	2	1	2	2			
P <sub>2</sub>			3	0	2	9	0	2	6	0	0			
P <sub>3</sub>			2	1	1	2	2	2	0	1	1			
P <sub>4</sub>			0	0	2	4	3	3	4	3	1			

- The OS is in a *safe state* since  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  is a safe sequence.



# Request Resource Algorithm

- 1 If ( $\text{Need}[i] < \text{Request}[i]$ ) then *max-error*
- 2 While ( $\text{Available} < \text{Request}[i]$ ) do *wait*
- 3 Construct a new state by:
  - $\text{Available} = \text{Available} - \text{Request}[i]$
  - $\text{Allocation}[i] = \text{Allocation}[i] + \text{Request}[i]$
  - $\text{Need}[i] = \text{Need}[i] - \text{Request}[i]$
- 4 If (*new state is not safe*) then  
*restore and wait*

# Request Example 1

- At some time,  $P_1$  requests an additional 1 A instance and 2 C instances
  - i.e.  $\text{Request}[1] == (1, 0, 2)$
- Does this lead to a safe state?
  - $\text{Available} \geq \text{Request}[1]$  so continue
  - generate new state and test for safety



			<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
A	B	C		A B C	A B C	A B C
P <sub>0</sub>			0 1 0	7 5 3	2 3 0	7 4 3
P <sub>1</sub>			3 0 2	3 2 2	0 2 0	
P <sub>2</sub>			3 0 2	9 0 2	6 0 0	
P <sub>3</sub>			2 1 1	2 2 2	0 1 1	
P <sub>4</sub>			0 0 2	4 3 3	4 3 3	

- The OS is in a *safe state* since  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  is a safe sequence.

# Further Request Examples

- From this state,  $P_4$  requests a further (3,3,0)
  - cannot be granted, since insufficient resources
- Alternatively,  $P_0$  requests a further (0,2,0)
  - should not be granted since the resulting state is unsafe



# Weaknesses in the Banker's Algorithm

- Weaknesses in the Banker's Algorithm
  - Requires there be a fixed number of resource to allocate
  - Requires the population of processes to be fixed
  - Requires the banker to grant all requests within "finite time"
  - Requires that clients repay all loans within "finite time"
  - Requires processes to state maximum needs in advance

# 5. Deadlock Detection

- If there are no prevention or avoidance mechanisms in place, then deadlock may occur.
- Deadlock detection should return enough information so the OS can recover.
- How often should the detection algorithm be executed?

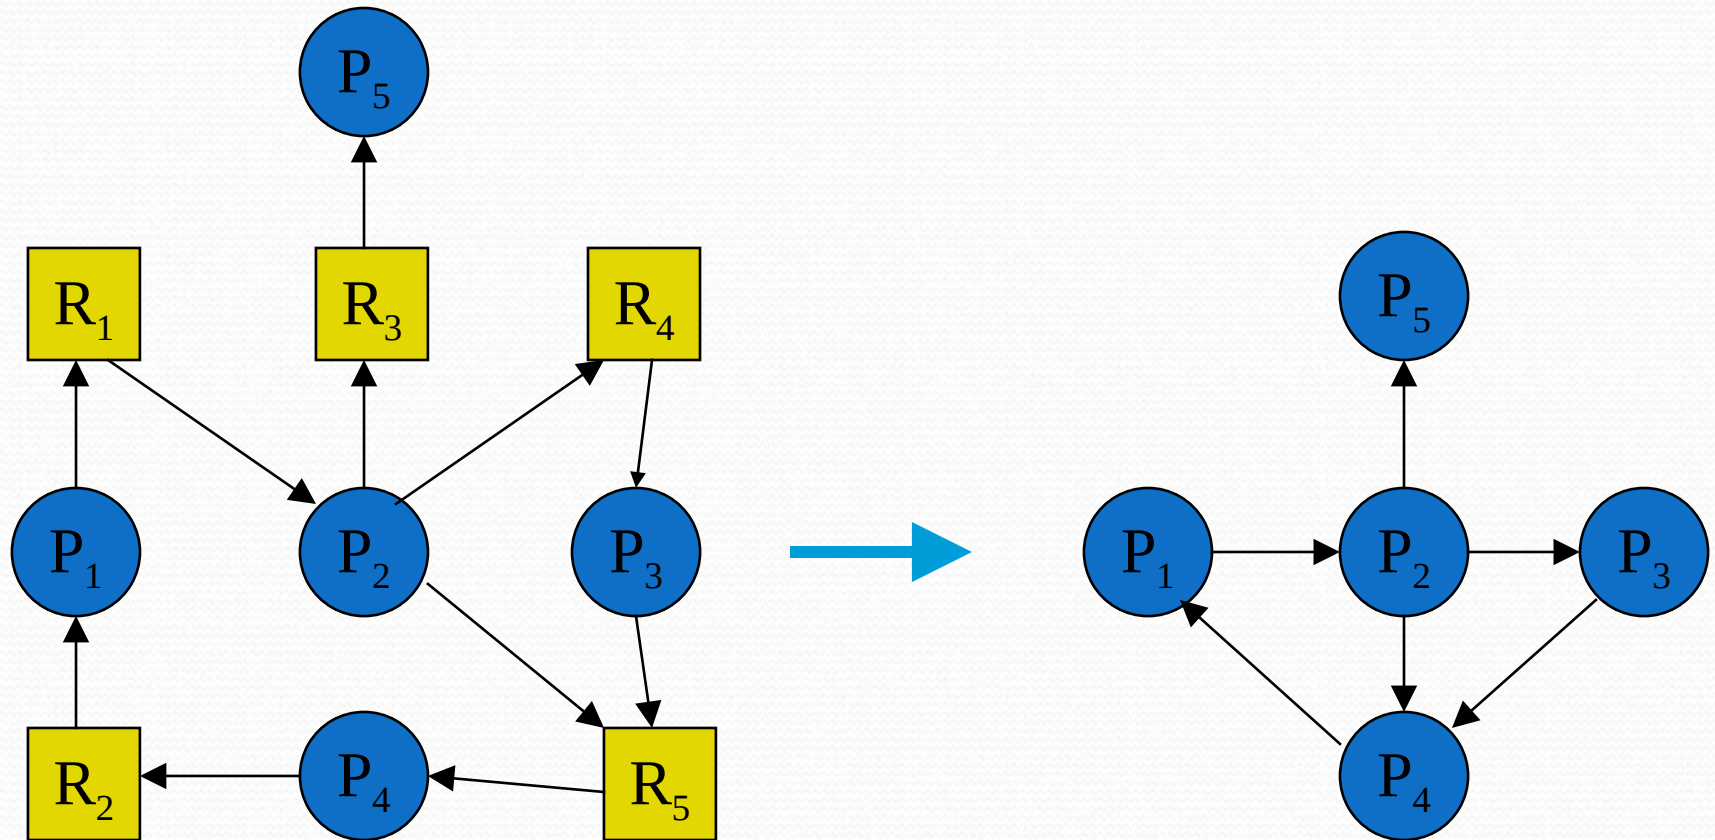


# 5.1. Wait-for Graph

- Assume that each resource has only one instance.
- Create a wait-for graph by removing the resource types nodes from a resource allocation graph.
- Deadlock exists if and only if the wait-for graph contains a cycle.

# Examp le

Fig. 7.7, p.225





## 5.2. Banker's Algorithm Variation

- If a resource type can have multiple instances, then an algorithm very similar to the banker's algorithm can be used.
- The algorithm investigates every possible allocation sequence for the processes that remain to be completed.

# Detection Algorithm

- 1 Vector Copy:  $Work := Available$ ;  $Finish := false$
- 2 Find  $i$  such that  $P_i$  hasn't finished but could:  
 $Finish[i] == false$   
 $Request[i] \leq Work$   
If no suitable  $i$ , go to step 4.
- 3 Assume  $P_i$  completes:  
 $Work := Work + Allocation[i]$   
 $Finish[i] := true$   
go to step 2
- 4 If  $Finish[i] == false$  then  $P_i$  is deadlocked



# Example

## 1

- | <u>Resource Type</u> | <u>Instances</u> |
|----------------------|------------------|
| A                    | 7                |
| B                    | 2                |
| C                    | 6                |

*continued*

<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
A	B	C	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	0	0	0	0	0
P <sub>1</sub>	2	0	0	0	2			
P <sub>2</sub>	3	0	0	0	0			
P <sub>3</sub>	2	1	1	0	0			
P <sub>4</sub>	0	0	2	0	2			

- The OS is *not* in a deadlocked state since  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  is a safe sequence.



# Example 2

- Change  $P_2$  to request 1 C instance

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	1			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- ❖ The OS *is* deadlocked.

# 6. Deadlock Recovery

- Tell the operator
- System-based recovery:
  - abort one or more processes in the circular wait
  - preempt resources in one or more deadlocked processes



# 6.1. Process Termination

- Abort all deadlocked processes, or
- Abort one process at a time until the deadlocked cycle disappears
  - not always easy to abort a process
  - choice should be based on *minimum cost*

# 6.1. Process Termination

The term *minimum cost* depends on various factors:

- What is the priority of the process
- How long the process has computed and how much longer the process will compute before completing its designated task
- How many and what types of resources the process has used.



# 6.1. Process Termination

The term *minimum cost* depends on various factors:

- How many more resources the process needs in order to complete
- How many processes will need to be terminated.
- Whether the process is interactive or batch

# 6.2. Resource Preemption

- Issues:
  - how to select a resource (e.g. by using minimum cost)
  - how to rollback the process which has just lost its resources
  - avoiding process starvation