

In a multiprogramming environment, several processes may compete for a finite number of resources. A process requests resources; and if the resources are not available at that time, the process enters a waiting state. Sometimes, a waiting process is never again able to change state, because the resources it has requested are held by other waiting processes.

**Deadlock** is defined as the permanent blocking of a set of processes that **compete** for system resources.

## 1. System model

- A system consists of a finite number of resources to be distributed among a number of competing processes.
- The resources are partitioned into several types, each consisting of some number of identical instances.
- **Reusable**: something that can be safely used by one process at a time and is not consumed by that use. Processes obtain resources that they later release for reuse by others (processors, memory, files, devices, databases, and semaphores).
- **Consumable**: these can be created and destroyed. When a resource is acquired by a process, the resource ceases to exist (interrupts, signals, messages, and information in I/O buffers).
- A **preemptable resource** is one that can be taken away from the process owning it with no ill effects. Memory (also CPU) is an example of a preemptable resource.
- A **non-preemptable resource**, in contrast, is one that cannot be taken away from its current owner without causing the computation to fail (printer, CD-RW, floppy disk).
- In general, deadlocks occur when sharing reusable and nonpreemptable resources. Potential deadlocks that involve preemptable resources can usually be resolved by reallocating resources from one process to another.
- A process must request a resource before using it and must release the resource after using it.
  - A process may request as many resources as it requires to carry out its designated task.
  - Obviously, the number of resources requested may not exceed the total number of resources available in the system.
- Under the normal mode of operation, a process may utilize a resource in only the following sequence:
  1. **Request**. If the request cannot be granted immediately, then the requesting process must wait until it can acquire the resource.
  2. **Use**. The process can operate on the resource.
  3. **Release**. The process releases the resource.
- The request and release of resources are system calls. Examples are the request ( ) and release ( ) device, open( ) and close( ) file, and malloc ( ) and free ( ) memory system calls.
- Request and release of resources that are not managed by the OS can be accomplished through the wait ( ) and signal ( ) operations on semaphores or through acquisition and release of a mutex(tools of synchronization) lock.
- A system table records whether each resource is free or allocated; for each resource that is allocated, the table also records the process to which it is allocated.
- If a process requests a resource that is currently allocated to another process, it can be added to a queue of processes waiting for this resource.
- A process whose resource request has just been denied will normally sit in a tight loop requesting the resource, then sleeping, then trying again.

- One possible way of allowing user management of resources is to associate a semaphore with each resource. Mutexes can be used equally well.
- The three steps listed above are then implemented as a down on the semaphore to acquire the resource, using the resource, and finally an up on the resource to release it.
- These steps are shown in Figure.

```
typedef int semaphore;
semaphore resource_1;
```

```
void process_A(void) {
    down(&resource_1);
    use_resource_1( );
    up(&resource_1);
}
```

(a)

```
typedef int semaphore;
semaphore resource_1;
semaphore resource_2;
```

```
void process_A(void) {
    down(&resource_1);
    down(&resource_2);
    use_both_resources( );
    up(&resource_2);
    up(&resource_1);
}
```

(b)

**Figure:** Using a semaphore to protect resources. (a) One resource. (b) Two resources.

- A set of processes is in a deadlock state when every process in the set is waiting for an event that can be caused only by another process in the set.
  - Because all the processes are waiting, none of them will ever cause any of the events that could wake up any of the other members of the set, and all the processes continue to wait forever.
  - None of the processes can run, none of them can release any resources, and none of them can be awakened.
  - This result holds for any kind of resource, including both hardware and software.
- To illustrate a deadlock state, consider a system with three CD RW drives.
  - Suppose each of three processes holds one of these CD RW drives.
  - If each process now requests another drive, the three processes will be in a deadlock state.
  - Each is waiting for the event "CD RW is released," which can be caused only by one of the other waiting processes.
- Deadlocks can occur in a variety of situations besides requesting dedicated I/O devices. In a database system, for example, a program may have to lock several records it is using, to avoid race conditions.
  - If process A locks record  $R_1$  and process B locks record  $R_2$ , and then each process tries to lock the other one's record, we also have a deadlock.

<pre> typedef int semaphore;     semaphore resource_1;     semaphore resource_2;      void process_A(void) {         down(&amp;resource_1);         down(&amp;resource_2);         use_both_resources( );         up(&amp;resource_2);         up(&amp;resource_1);     }      void process_B(void) {         down(&amp;resource_1);         down(&amp;resource_2);         use_both_resources( );         up(&amp;resource_2);         up(&amp;resource_1);     } </pre> <p style="text-align: center;">(a)</p>	<pre> semaphore resource_1; semaphore resource_2;  void process_A(void) {     down(&amp;resource_1);     down(&amp;resource_2);     use_both_resources( );     up(&amp;resource_2);     up(&amp;resource_1); }  void process_B(void) {     down(&amp;resource_2);     down(&amp;resource_1);     use_both_resources( );     up(&amp;resource_1);     up(&amp;resource_2); } </pre> <p style="text-align: center;">(b)</p>
--	---

**Figure:** (a) Deadlock-free code. (b) Code with a potential deadlock.

- Deadlocks can occur on hardware resources or on software resources.
- Unlike other problems in multiprogramming systems, there is no efficient solution to the deadlock problem in the general case.
- A programmer who is developing multithreaded applications must pay particular attention to this problem. Multithreaded programs are good candidates for deadlock because multiple threads can compete for shared resources.

## 2. Necessary Conditions

- A deadlock situation can arise if the following four conditions hold simultaneously in a system:
  - **Mutual exclusion.** At least one resource must be held in a nonsharable mode;
    - That is, only one process at a time can use the resource.
    - If another process requests that resource, the requesting process must be delayed until the resource has been released.
  - **Hold and wait.** A process must be holding at least one resource and waiting to acquire additional resources that are currently being held by other processes.
  - **No preemption.** Resources cannot be preempted; that is, a resource can be released only voluntarily by the process holding it, after that process has completed its task.
  - **Circular wait.** A set  $P_0, P_1, \dots, P_n$  of waiting processes must exist such that
    - $P_0$  is waiting for a resource held by  $P_1$ ,
    - $P_1$  is waiting for a resource held by  $P_2$ ,
    - .
    - .

- .
- $P_{n-1}$  is waiting for a resource held by  $P_n$ ,
- $P_n$  waiting for a resource held by  $P_0$ .

There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain.

We emphasize that **all four conditions must hold for a deadlock** to occur.

### 3. Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

$V$  is partitioned into two types:

$P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.

$R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.

Request edge – directed edge  $P_i \longrightarrow R_j$

Assignment edge – directed edge  $R_j \longrightarrow P_i$

Process



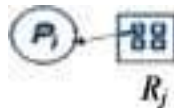
Resource Type with 4 instances



$P_i$  requests instance of



$P_i$  is holding an instance of  $R_j$



The resource-allocation graph is shown in Figure which depicts the following situation.

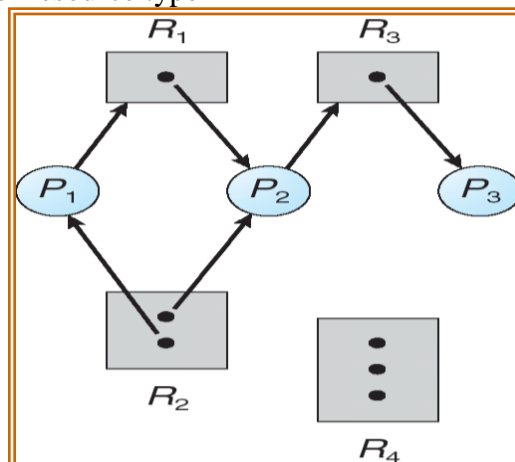
The sets  $P$ ,  $R$ , and  $E$ :

$P = \{P_1, P_2, P_3\}$

$R = \{R_1, R_2, R_3, R_4\}$

Resource instances:

- One instance of resource type  $R_1$
- Two instances of resource type  $R_2$
- One instance of resource type  $R_3$
- Three instances of resource type  $R_4$



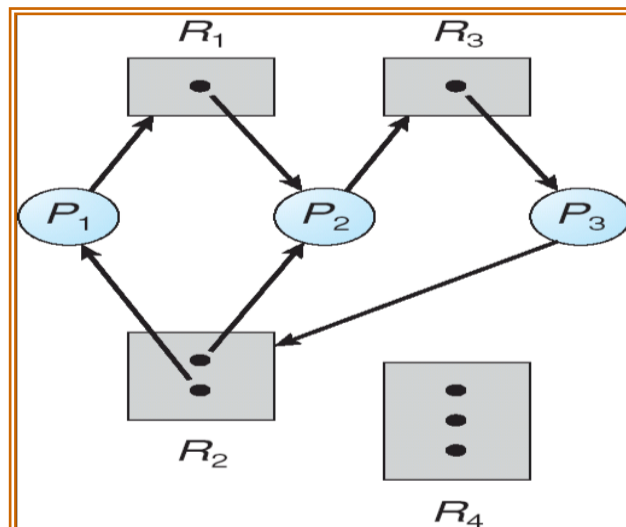
Resource-allocation graph

Process status:

- Process  $P_1$  is holding an instance of resource type  $R_2$ , and is waiting for an instance of resource type  $R_1$
- Process  $P_2$  is holding an instance of  $R_1$  and  $R_2$  and is waiting for an instance of resource type  $R_3$
- $P_3$  is holding an instance of  $R_3$

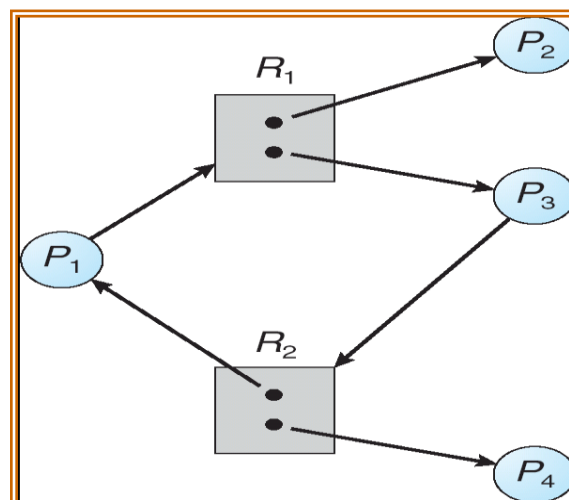
Suppose that process  $P_3$  requests an instance of resource type  $R_2$ . Since no resource instance is currently available, a request edge  $P_3 \rightarrow R_2$  is added to the graph (Figure 7.2). At this point, two minimal cycles exist in the system:

$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$   
 $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$



However, there is no deadlock. Observe that process  $P_2$  may release its instance of resource type  $R_2$ . That resource can then be allocated to  $P_3$ , breaking the cycle.

In summary, if a resource-allocation graph does not have a cycle, then the system is *not* in a deadlock state. On the other hand, if there is a cycle, then the system may or may not be in a deadlock state. This observation is important when we deal with the deadlock problem.



Resource-allocation graph with a cycle but no deadlock.

#### 4. Methods for Handling Deadlocks

Generally speaking, we can deal with the deadlock problem in one of three ways:

1. We can use a protocol to **prevent** or **avoid** deadlocks, ensuring that the system will never enter a deadlock state.
    - To ensure that deadlocks never occur, the system can use either a deadlock-prevention or a deadlock-avoidance scheme.
    - Deadlock prevention provides a set of methods for ensuring that at least one of the necessary conditions cannot hold.
    - Design a system in such a way that the possibility of deadlock is excluded a priori (compile-time/statically, by design).
    - Deadlock avoidance requires that the OS be given in advance additional information concerning which resources a process will request and use during its lifetime. With this additional knowledge, it can decide for each request whether or not the process should wait.
    - Make a decision dynamically checking whether the request will, if granted, potentially lead to a deadlock or not. (run-time/dynamically, before it happens).
  2. We can allow the system to enter a deadlock state, **detect** it, and **recover**.
    - If a system does not employ either a deadlock-prevention or a deadlock-avoidance algorithm, then a deadlock situation may arise.
    - In this environment, the system can provide an algorithm that examines the **state** of the system to **determine** whether a deadlock has occurred and an algorithm to **recover** from the deadlock.
    - Let deadlocks occur, detect them, and take action (run-time/dynamically, after it happens)
  3. We can **ignore** (The Ostrich Algorithm; maybe if you ignore it, it will ignore you) the problem altogether and pretend that deadlocks never occur in the system.
    - If a system neither ensures that a deadlock will never occur nor provides a mechanism for deadlock detection and recovery, then we may arrive at a situation where the system is in a deadlocked state yet has no way of recognizing what has happened.
    - Eventually, the system will stop functioning and will need to be restarted manually.
    - In many systems, deadlocks occur infrequently (say, once per year); thus, this method is cheaper than the prevention, avoidance, or detection and recovery methods.
- Most OSs potentially suffer from deadlocks that are not even detected. Process table slots are finite resources. If a fork fails because the table is full, a reasonable approach for the program doing the fork is to wait a random time and try again.
  - The maximum number of open files is similarly restricted by the size of the i-node table, so a similar problem occurs when it fills up. Swap space on the disk is another limited resource. In fact, almost every table in the OS represents a finite resource.
  - The third solution is the one used by most OSs, including UNIX and Windows; it is then up to the application developer to write programs that handle deadlocks.
  - If deadlocks could be eliminated for free, there would not be much discussion. The problem is that the price is high, mostly in terms of putting inconvenient restrictions on processes.
  - Thus we are faced with an unpleasant trade-off between convenience and correctness. Under these conditions, general solutions are hard to find.

## 4.1 Deadlock Prevention

- Having seen that deadlock avoidance is essentially impossible, because it requires information about future requests, which is not known, how do real systems avoid deadlock?
- If we can ensure that at least one of the four following conditions is never satisfied, then deadlocks will be structurally impossible.
- The various approaches to deadlock prevention are summarized in Figure.

Condition	Approach
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away
Circular wait	Order resources numerically

**Figure:** Summary of approaches, to deadlock prevention.

### 4.1.1 Mutual Exclusion

- **Attacking the Mutual Exclusion Condition;** can a given resource be assigned to more than one process at once? Systems with only simultaneously shared resources cannot deadlock!
- The mutual-exclusion condition must hold for non-sharable resources (i.e., a printer).
- Shareable resources, in contrast, do not require mutually exclusive access and thus cannot be involved in a deadlock (i.e. read-only files). A process never needs to wait for a shareable resource.
- In general, however, we cannot prevent deadlocks by denying the mutual-exclusion condition, because some resources are intrinsically non-sharable.

### 4.1.2 Hold and Wait

- **Attacking the Hold and Wait Condition;** can a process hold a resource and ask for another? Can we require processes to request all resources at once? Most processes do not statically know about the resources they need.
- To ensure that the hold-and-wait condition never occurs in the system, we must guarantee that, whenever a process requests a resource, it does not hold any other resources.
- One protocol that can be used requires each process to request and be allocated all its resources before it begins execution.
- An alternative protocol allows a process to request resources only when it has none. A process may request some resources and use them. Before it can request any additional resources, however, it must release all the resources that it is currently allocated.
- To illustrate the difference between these two protocols, we consider a process that copies data from a DVD drive to a file on disk, sorts the file, and then prints the results to a printer.
- Both these protocols have two main disadvantages.
  - First, resource utilization may be low, since resources may be allocated but unused for a long period.
  - Second, starvation is possible. A process that needs several popular resources may have to wait indefinitely, because at least one of the resources that it needs is always allocated to some other process.

### 4.1.3 No Preemption

- **Attacking the No Preemption Condition;** Can resources be preempted? If a process' requests (holding certain resources) is denied, that process must release its unused resources and request them again, together with the additional resource.

- The third necessary condition for deadlocks is that there be no preemption of resources that have already been allocated.
- To ensure that this condition does not hold, we can use the following protocol.
  - If a process is holding some resources and requests another resource that cannot be immediately allocated to it (that is, the process must wait), then all resources currently being held are preempted.
  - The preempted resources are added to the list of resources for which the process is waiting. The process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

#### 4.1.4 Circular Wait

- **Attacking the Circular Wait Condition;** Can circular waits exist? Order resources by an index:  $R_1, R_2, \dots$ ; requires that resources are always requested in order.
- One way to ensure that this condition never holds is to impose a total ordering of all resource types and to require that each process requests resources in an increasing order of enumeration.
- Assign to each resource type a unique integer number, which allows us to compare two resources and to determine whether one precedes another in our ordering.
- Each process can request resources only in an increasing order of enumeration.
- If these two protocols are used, then the circular-wait condition cannot hold.

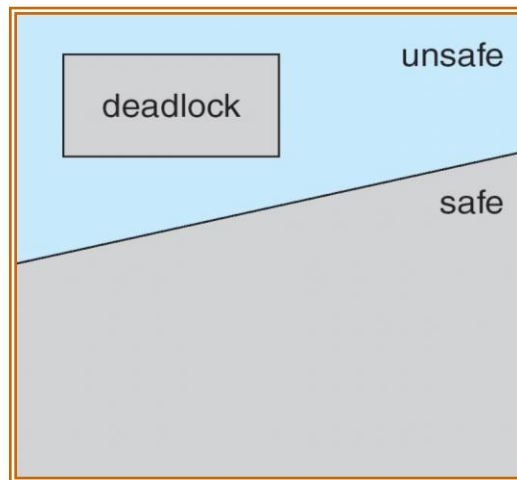
#### 4.2 Deadlock Avoidance

- Deadlock-prevention algorithms ensure that at least one of the necessary conditions for deadlock cannot occur and hence that deadlocks cannot hold.
- Possible side effects of preventing deadlocks are low device utilization and reduced system throughput.
- An alternative method for avoiding deadlocks is to require additional information about how resources are to be requested.
  - For example, in a system with one tape drive and one printer, the system might need to know that process P will request first the tape drive and then the printer before releasing both resources, whereas process Q will request first the printer and then the tape drive.
  - With this knowledge of the complete sequence of requests and releases for each process, the system can decide for each request whether or not the process should wait in order to avoid a possible future deadlock.
- Each request requires that in making this decision the system consider
  - the resources currently available,
  - the resources currently allocated to each process,
  - the future requests and releases of each process.
- The simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need.
- Given this a priori information, it is possible to construct an algorithm that ensures that the system will never enter a deadlocked state.
- A deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that a circular-wait condition can never exist.
- The resource-allocation state is defined by the number of available and allocated resources and the maximum demands of the processes.



#### 4. Safe State

- A state is **safe** if the system can allocate resources to each process (up to its maximum) in some order and still avoid a deadlock.
- More formally, a system is in a safe state only if there exists a **safe sequence**.
  - A sequence of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is a safe sequence for the current allocation state if, for each  $P_i$  the resource requests that  $P_i$  can still make can be satisfied by the currently available resources plus the resources held by all  $P_j$ , with  $j < i$ .
  - In this situation, if the resources that  $P_i$  needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
  - When they have finished,  $P_i$  can obtain all of its needed resources, complete its designated task, return its allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.
  - If no such sequence exists, then the system state is said to be **unsafe**.
- A safe state is not a deadlocked state. Conversely, a deadlocked state is an unsafe state. Not all unsafe states are deadlocks, however Figure.



**Figure:** Safe, unsafe, and deadlock state spaces.

- An unsafe state may lead to a deadlock.
    - As long as the state is safe, the OS can avoid unsafe (and deadlocked) states.
    - In an unsafe state, the OS cannot prevent processes from requesting resources such that a deadlock occurs: The behavior of the processes controls unsafe states.
    - The difference between a safe state and an unsafe state is that from a safe state the system can guarantee that all processes will finish; from an unsafe state, no such guarantee can be given.
- Example:**
- To illustrate, we consider a system with 12 magnetic tape drives and three processes:  $P_0, P_1$ , and  $P_2$ .
    - Process  $P_0$  requires 10 tape drives, and holds 5 tape drives
    - Process  $P_1$  may need as many as 4 tape drives, and holds 2 tape drive
    - Process  $P_2$  may need up to 9 tape drives, and holds 2 tape drives
    - Thus, there are 3 free tape drives.

$P_i$	Maximum Needs	Current Needs
$P_0$	10	5
$P_1$	4	2
$P_2$	9	7

- At time  $t_0$ , the system is in a safe state. The sequence  $\langle P_1, P_0, P_2 \rangle$  satisfies the safety condition.
- A system can go from a safe state to an unsafe state. Suppose that, at time  $t_1$ , process  $P_2$  requests and is allocated one more tape drive. The system is no longer in a safe state.
  - At this point, only process  $P_1$  can be allocated all its tape drives.
  - When it returns them, the system will have only 4 available tape drives.
  - Since process  $P_0$  is allocated 5 tape drives but has a maximum of 10, it may request 5 more tape drives. Since they are unavailable, process  $P_0$  must wait.
  - Similarly, process  $P_2$  may request an additional 6 tape drives and have to wait, resulting in a deadlock.
- Our mistake was in granting the request from process  $P_2$  for one more tape drive.

#### Example:

- In Figure we have a state in which
  - A total of 10 instances of the resource exist, so with 7 resources already allocated, there are 3 still free.
  - A has 3 instances of the resource but may need as many as 9 eventually.
  - B currently has 2 and may need 4 altogether, later.
  - Similarly, C also has 2 but may need an additional 5.
  - The state of Fig. upper is safe because there exists a sequence of allocations that allows all processes to complete. By careful scheduling, can avoid deadlock.
  - Now suppose, this time A requests and gets another resource, giving Figure lower. Eventually, B completes. At this point we are stuck.
  - We only have four instances of the resource free, and each of the active processes needs five. There is no sequence that guarantees completion. A's request should not have been granted.
  - But, an unsafe state is not a deadlocked state. It is possible that A might release a resource before asking for any more, allowing C to complete and avoiding deadlock altogether.

Has Max			Has Max			Has Max			Has Max			Has Max		
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	-	B	0	-	B	0	-
C	2	7	C	2	7	C	2	7	C	7	7	C	0	-
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7		
(a)			(b)			(c)			(d)			(e)		

Has Max			Has Max			Has Max			Has Max		
A	3	9	A	4	9	A	4	9	A	4	9
B	2	4	B	2	4	B	4	4	B	—	—
C	2	7	C	2	7	C	2	7	C	2	7
Free: 3			Free: 2			Free: 0			Free: 4		
(a)			(b)			(c)			(d)		

**Figure:** Demonstration that the state in is safe (Upper), is not safe (Lower).

- Given the concept of a safe state, we can define avoidance algorithms that ensure that the system will never deadlock.
  - The idea is simply to ensure that the system will always remain in a safe state.
  - Initially, the system is in a safe state.
  - Whenever a process requests a resource that is currently available, the system must decide whether the resource can be allocated immediately or whether the process must wait.
  - The request is granted only if the allocation leaves the system in a safe state.
- In this scheme, if a process requests a resource that is currently available, it may still have to wait. Thus, resource utilization may be lower than it would otherwise be.

## 5. Banker's Algorithm for resource allocation

### Some conditions of Banker's Algorithm

- Multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

### Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- Max:**  $n \times m$  matrix. If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- Allocation:**  $n \times m$  matrix. If  $Allocation[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- Need:**  $n \times m$  matrix. If  $Need[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j].$$

### Safety Algorithm

- Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively. Initialize:  
 $Work = Available$   
 $Finish[i] = false$  for  $i = 1, 3, \dots, n$ .
- Find and  $i$  such that both:
  - $Finish[i] = false$
  - $Need_i \leq Work$
- If no such  $i$  exists, go to step 4.
- $Work = Work + Allocation_i$   
 $Finish[i] = true$   
 go to step 2.
- If  $Finish[i] == true$  for all  $i$ , then the system is in a safe state.

### Example of Banker's Algorithm

5 processes  $P_0$  through  $P_4$ ; 3 resource types A (10 instances), B (5 instances), and C (7 instances).

Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

- The content of the matrix. Need is defined to be  $\text{Max} - \text{Allocation}$ .

	<u>Need</u>
	A B C
$P_0$	7 4 3
$P_1$	1 2 2
$P_2$	6 0 0
$P_3$	0 1 1
$P_4$	4 3 1

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

Now  $P_1$  request (1, 0, 2)

- Check that  $\text{Request} \leq \text{Available}$  (that is,  $(1, 0, 2) \leq (3, 3, 2) \Rightarrow \text{true}$ ).

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 1	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

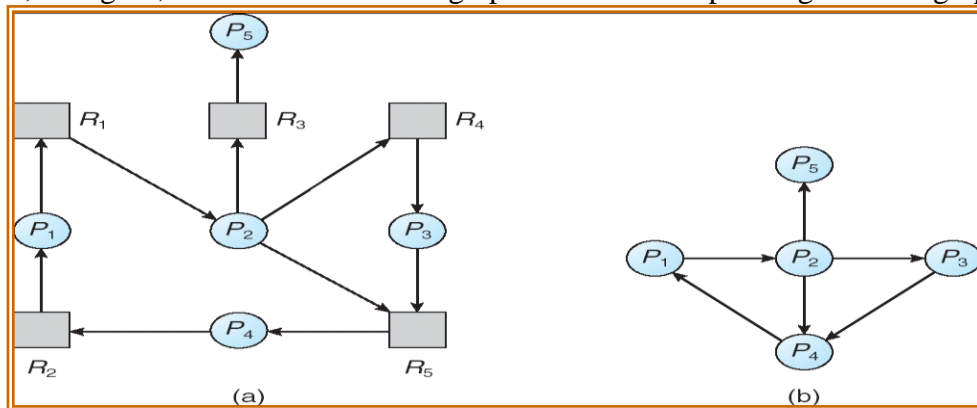
- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement.
- Can request for (3, 3, 0) by  $P_4$  be granted?
- Can request for (0, 2, 0) by  $P_0$  be granted?

## 6. Deadlock Detection

- If a system does not employ either a deadlock-prevention or a deadlock-avoidance algorithm then a deadlock situation may occur.
- In this environment, the system must provide:
  - An algorithm that examines the state of the system to determine whether a deadlock has occurred.
  - An algorithm to recover from the deadlock.

### 5.1 Single Instance of Each Resource Type

- If all resources have only a single instance, then we can define a deadlock-detection algorithm that uses a variant of the resource-allocation graph, called a wait-for graph.
- This graph is obtained from the resource-allocation graph by removing the resource nodes and collapsing the appropriate edges.
- More precisely, an edge from  $P_i$  to  $P_j$  in a wait-for graph implies that process  $P_i$  is waiting for process  $P_j$  to release a resource that  $P_i$  needs.
- For example, in Figure, a resource-allocation graph and the corresponding wait-for graph are



presented.

**Figure 7.11:** (a) Resource-allocation graph. (b) Corresponding wait-for graph

- As before, a deadlock exists in the system if and only if the wait-for graph contains a cycle.
- To detect deadlocks, the system needs to maintain the wait-for graph and periodically invoke an algorithm that searches for a cycle in the graph.
- If this graph contains one or more cycles (knots), a deadlock exists. Any process that is part of a cycle is deadlocked. If no cycles exist, the system is not deadlocked.

### 5.2 Several Instances of a Resource Type

- The wait-for graph scheme is not applicable to a resource-allocation system with multiple instances of each resource type.
- We turn now to a deadlock-detection algorithm that is applicable to such a system. The algorithm employs several time-varying data structures:
  - **Available.** A vector of length  $m$  indicates the number of available resources of each type.
  - **Allocation.** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
  - **Request.** An  $n \times m$  matrix indicates the current request of each process.
    - If  $\text{Request}[i][j]$  equals  $k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .
- The detection algorithm described here simply investigates every possible allocation sequence for the processes that remain to be completed.

1. Let work and finish be vectors of length m and n, respectively. Initialize work=available. For  $i=0,1,\dots,n-1$ , if  $\text{Allocation}_i \neq 0$ , then  $\text{finish}[i]=\text{false}$ ; otherwise,  $\text{finish}[i]=\text{true}$ .
2. Find an index i such that both
  - a.  $\text{finish}[i]=\text{false}$
  - b.  $\text{Request}_i \leq \text{work}$
 If no such i exists, go to step 4.
3.  $\text{Work}=\text{work}+\text{allocation}$   
 $\text{finish}[i]=\text{true}$   
 Go to step 2.
4. If  $\text{finish}[i]=\text{false}$ , for some i,  $0 \leq i \leq n$ , then the system is in a deadlocked state. Moreover, if  $\text{finish}[i]=\text{false}$ , then process  $P_i$  is deadlocked.  
**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.**

- Consider a system with five processes  $P_0$  through  $P_4$  and three resource types A, B, and C.
  - Resource type A has seven instances,
  - Resource type B has two instances,
  - Resource type C has six instances.
- Suppose that, at time  $T_0$ , we have the following resource-allocation state:

	Allocation	Request	Available
$P_i$	<i>ABC</i>	<i>ABC</i>	<i>ABC</i>
$P_0$	010	000	000
$P_1$	200	202	
$P_2$	303	000	
$P_3$	211	100	
$P_4$	002	002	

- If the algorithm is executed, it will be found that the sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  result in  $\text{finish}[i]=\text{true}$  for all i.
- Suppose now that process  $P_2$  makes one additional request for an instance of type C.
- Although we can reclaim the resources held by process  $P_0$ , the number of available resources is not sufficient to fulfill the requests of the other processes. Thus, a deadlock exists, consisting of processes  $P_1, P_2, P_3$ , and, and  $P_4$ .

### Recovery from Deadlock

- Process Termination
- Resource Preemption